

# Lecture 15 — Learning Theory (Part 2 of 2)

Alex Schwing and Matus Telgarsky

March 6, 2018

## Schedule for today.

- ▶ Midterm announcements.
- ▶ Overfitting/generalization: reminders.
- ▶ Overfitting/generalization: intuition and basics.
- ▶ Overfitting/generalization: some asides.

### **Learning theory reading:**

see my learning theory course's resources (click on <http://mjt.cs.illinois.edu/courses/mlt-f17/> or google "matus uiuc mlt-f17").

## Midterm announcements.

- ▶ **Location (all ECEB):** your netid determines your room:
  - ▶ **1002** (this room): aa18 - ryang28.
  - ▶ **1013:** sabag2 - xunlin2.
  - ▶ **1015:** xyu69 - zzhou51.
- ▶ **Time:** start time 6pm, duration 90 minutes.
- ▶ **Notes:** can bring one standard size (“US letter”) sheet of notes, front and back, **handwritten**.
- ▶ **Review lecture and materials:** wait until Thursday.

## Learning Theory (part 2 of 2) – Overfitting/Generalization.

- ▶ **Overfitting:** better performance on past data than future data.
- ▶ **Generalizing:** similar performance on past and future data.

# Learning Theory (part 2 of 2) – Overfitting/Generalization.

- ▶ **Overfitting:** better performance on past data than future data.
- ▶ **Generalizing:** similar performance on past and future data.

## Models

Reasoning about this requires a **model** linking past and future.

- ▶ **Statistical learning theory:** past and future examples drawn IID from a common distribution.
- ▶ **Online learning:** adversary constructs new examples.

## Overfitting by example.

Linear or polynomial least squares?


**Truth:**  $y = 0 \cdot x + \xi$ ,  $\xi \sim \text{Gaussian}$ .

# Overfitting by example.

Linear or polynomial least squares?

**Truth:**  $y = 0 \cdot x + \xi$ ,  $\xi \sim \text{Gaussian}$ .

**Red:** seen data. **Green:** unseen data.




# Overfitting by example.

Linear or polynomial least squares?

**Truth:**  $y = 0 \cdot x + \xi$ ,  $\xi \sim \text{Gaussian}$ .

**Red:** seen data. **Green:** unseen data.






# Overfitting by example.

Linear or polynomial least squares?

**Truth:**  $y = 0 \cdot x + \xi$ ,  $\xi \sim \text{Gaussian}$ .

**Red:** seen data. **Green:** unseen data.




# Overfitting by example.

Linear or polynomial least squares?

**Truth:**  $y = 0 \cdot x + \xi$ ,  $\xi \sim \text{Gaussian}$ .

**Red:** seen data. **Green:** unseen data.




# Overfitting by example.

Linear or polynomial least squares?

**Truth:**  $y = 0 \cdot x + \xi$ ,  $\xi \sim \text{Gaussian}$ .

**Red:** seen data. **Green:** unseen data.




# Overfitting by example.

Linear or polynomial least squares?

**Truth:**  $y = 0 \cdot x + \xi$ ,  $\xi \sim \text{Gaussian}$ .

**Red:** seen data. **Green:** unseen data.




# Overfitting by example.

Linear or polynomial least squares?

**Truth:**  $y = 0 \cdot x + \xi$ ,  $\xi \sim \text{Gaussian}$ .

**Red:** seen data. **Green:** unseen data.




# Overfitting by example.

Linear or polynomial least squares?

**Truth:**  $y = 0 \cdot x + \xi$ ,  $\xi \sim \text{Gaussian}$ .

**Red:** seen data. **Green:** unseen data.




# Overfitting by example.

Linear or polynomial least squares?

**Truth:**  $y = 0 \cdot x + \xi$ ,  $\xi \sim \text{Gaussian}$ .

**Red:** seen data. **Green:** unseen data.




# Overfitting by example.

Linear or polynomial least squares?

**Truth:**  $y = 0 \cdot x + \xi$ ,  $\xi \sim \text{Gaussian}$ .

**Red:** seen data. **Green:** unseen data.






# Overfitting by example.

Linear or polynomial least squares?

**Truth:**  $y = 0 \cdot x + \xi$ ,  $\xi \sim \text{Gaussian}$ .

**Red:** seen data. **Green:** unseen data.




# Overfitting by example.

Linear or polynomial least squares?

**Truth:**  $y = 0 \cdot x + \xi$ ,  $\xi \sim \text{Gaussian}$ .

**Red:** seen data. **Green:** unseen data.

$\hat{R}_1 = 0.00692304, R_1 = 0.00897457, \hat{R}_{21} = 0.0026458, R_{21} = 1.24795e + 07$




# Overfitting by example.

Linear or polynomial least squares?

**Truth:**  $y = 0 \cdot x + \xi$ ,  $\xi \sim \text{Gaussian}$ .

**Red:** seen data. **Green:** unseen data.




# Overfitting by example.

Linear or polynomial least squares?

**Truth:**  $y = 0 \cdot x + \xi$ ,  $\xi \sim \text{Gaussian}$ .

**Red:** seen data. **Green:** unseen data.

$\hat{R}_1 = 0.00692304, R_1 = 0.00897457, \hat{R}_{34} = 0.00259486, R_{34} = 3.97751e + 10$




# Overfitting by example.

Linear or polynomial least squares?

**Truth:**  $y = 0 \cdot x + \xi$ ,  $\xi \sim \text{Gaussian}$ .

**Red:** seen data. **Green:** unseen data.

$\hat{R}_1 = 0.00692304, R_1 = 0.00897457, \hat{R}_{38} = 0.00242094, R_{38} = 1.02285e + 13$



# Overfitting by example.




Figure 1: Fitting a degree 38 polynomial.

**Intuition:** More parameters  $\implies$  more overfitting.

**Can reduce overfitting** with model choice and regularization (similar...).

**What learning theory gives us:** concrete relationships.

Aside/review: least squares code for this plot!

```
# [ . . . ]

for s in [ 'tr', 'te', 'grid' ]:
    X[s][:, 0] = 1.0
    #X[s][:, 1] is random according to some distribution
    for j in range(2, n):
        X[s][:, j] = X[s][:, 1] * X[s][:, j - 1]

# [ . . . ]

for j in range(2, n):
    #better to use the black box N.linalg.lstsq
    w[j] = N.linalg.pinv(X['tr'][:, :j]) @ Y['tr']
    R[j] = dict( (s, N.linalg.norm(X[s][:, :j] @ w[j]
                                   - Y[s])**2 / 2 / n)
                 for s in [ 'tr', 'te' ] )
```

Binomials, random walks, coin tosses, classification.



# Binomials, random walks, coin tosses, classification.

Last lecture we saw *Hoeffding's Inequality*.

**Theorem** (Hoeffding's inequality). Suppose each draw from the distribution lies in the interval  $[a, b]$ . With probability at least  $1 - \delta$  over an iid draw of  $(z_i)_{i=1}^n$ ,

$$\mathbb{E}Z \leq \frac{1}{n} \sum_{i=1}^n z_i + (b - a) \sqrt{\frac{\ln(1/\delta)}{2n}}.$$

# Binomials, random walks, coin tosses, classification.

Last lecture we saw *Hoeffding's Inequality*.

**Theorem** (Hoeffding's inequality). Suppose each draw from the distribution lies in the interval  $[a, b]$ . With probability at least  $1 - \delta$  over an iid draw of  $(z_i)_{i=1}^n$ ,

$$\mathbb{E}Z \leq \frac{1}{n} \sum_{i=1}^n z_i + (b - a) \sqrt{\frac{\ln(1/\delta)}{2n}}.$$


What does this mean?

## Interpreting Hoeffding's inequality.

```
for i in range(k):  
    path = N.cumsum( N.random.randint(0, 2, n) * 2 - 1 )  
    plt.plot(path, color = 'black', alpha = 0.35)
```


## Interpreting Hoeffding's inequality.

```
for i in range(k):  
    path = N.cumsum( N.random.randint(0, 2, n) * 2 - 1 )  
    plt.plot(path, color = 'black', alpha = 0.35)
```



## Interpreting Hoeffding's inequality.


```
for i in range(k):  
    path = N.cumsum( N.random.randint(0, 2, n) * 2 - 1 )  
    plt.plot(path, color = 'black', alpha = 0.35)
```



Hoeffding says: for any **fixed**  $n$ , with probability  $\geq 1 - 1/\sqrt{e}$ ,  
$$\text{position} \leq \sqrt{n}.$$

## Hoeffding's inequality with more walks.


As paths are added, the upper bound **must** increase.



**Question:** how is the curve being rescaled?

## Hoeffding's inequality with more walks.


As paths are added, the upper bound **must** increase.



**Question:** how is the curve being rescaled?

## Hoeffding's inequality with more walks.

As paths are added, the upper bound **must** increase.




**Question:** how is the curve being rescaled?



## Hoeffding's inequality with more walks.

As paths are added, the upper bound **must** increase.



**Question:** how is the curve being rescaled?

## Hoeffding's inequality with more walks.


As paths are added, the upper bound **must** increase.



**Question:** how is the curve being rescaled?

## Hoeffding's inequality with more walks.


As paths are added, the upper bound **must** increase.



**Question:** how is the curve being rescaled?

## Hoeffding's inequality with more walks.


As paths are added, the upper bound **must** increase.



**Question:** how is the curve being rescaled?

## Hoeffding's inequality with more walks.


As paths are added, the upper bound **must** increase.



**Question:** how is the curve being rescaled?

## Hoeffding's inequality with more walks.

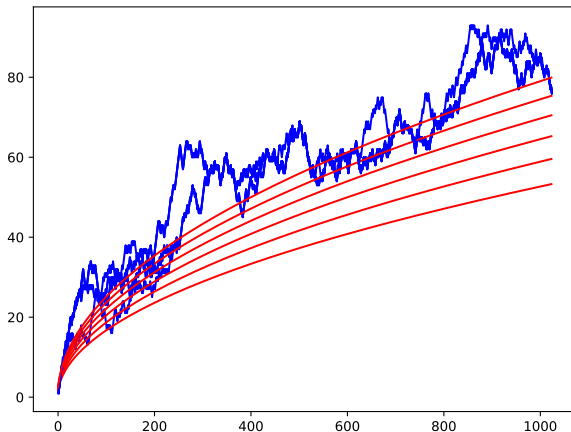
As paths are added, the upper bound **must** increase.



**Question:** how is the curve being rescaled?

## Hoeffding's inequality with more walks.

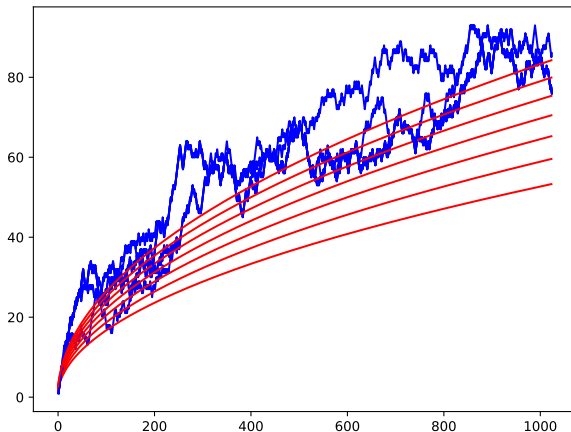
As paths are added, the upper bound **must** increase.



**Question:** how is the curve being rescaled?

## Hoefding's inequality with more walks.

As paths are added, the upper bound **must** increase.



**Question:** how is the curve being rescaled?



## Helpful tool: the union bound.

Given events  $(E_1, \dots, E_k)$ , the probability that *some*  $E_i$  occurs is

$$\Pr(E_1 \vee \dots \vee E_k) \leq \sum_i \Pr(E_i).$$

**(Intuition:** Venn diagram.)

## Helpful tool: the union bound.

Given events  $(E_1, \dots, E_k)$ , the probability that *some*  $E_i$  occurs is

$$\Pr(E_1 \vee \dots \vee E_k) \leq \sum_i \Pr(E_i).$$

**(Intuition:** Venn diagram.)

A useful consequence.

Probability that *no*  $E_i$  occurs satisfies

$$\Pr(\neg E_1 \wedge \dots \wedge \neg E_k) \geq 1 - \sum_i \Pr(E_i).$$

## Bounding *many* paths using Hoeffding.

**Theorem** (Hoeffding's inequality). Consider random variables  $(W_1, \dots, W_k)$  where  $W_j := \frac{1}{n} \sum_{i=1}^n Z_{j,i}$  with  $Z_{j,i} \in [a, b]$ , and  $Z_{j,i}$  are independent for fixed  $j$ , but may be *dependent* for fixed  $i$ . With probability at least  $1 - \delta$ ,

$$\max_{j \in [k]} (\mathbb{E}(W_j) - W_j) \leq (b - a) \sqrt{\frac{\ln(k) + \ln(1/\delta)}{2n}}.$$

## Bounding *many* paths using Hoeffding.

**Theorem** (Hoeffding's inequality). Consider random variables  $(W_1, \dots, W_k)$  where  $W_j := \frac{1}{n} \sum_{i=1}^n Z_{j,i}$  with  $Z_{j,i} \in [a, b]$ , and  $Z_{j,i}$  are independent for fixed  $j$ , but may be *dependent* for fixed  $i$ . With probability at least  $1 - \delta$ ,

$$\max_{j \in [k]} (\mathbb{E}(W_j) - W_j) \leq (b - a) \sqrt{\frac{\ln(k) + \ln(1/\delta)}{2n}}.$$


► **Proof.** Define  $\epsilon := \sqrt{\ln(k/\delta)/2n}$  and events

$$E_j := [\mathbb{E}(W_j) > W_j + \epsilon].$$

By Hoeffding,  $\Pr(E_j) \leq \delta/k$ , and by union bound

$$\Pr(\cap_j \neg E_j) \geq 1 - \sum_j \Pr(E_j) \geq 1 - \delta.$$


## Back to bounding random walks.



**Question:** how is the curve being rescaled?

**Answer:**  $\sqrt{\ln(\#\text{paths})}$ .


## Back to bounding random walks.



**Question:** how is the curve being rescaled?

**Answer:**  $\sqrt{\ln(\#\text{paths})}$ .


## Back to bounding random walks.



**Question:** how is the curve being rescaled?

**Answer:**  $\sqrt{\ln(\#\text{paths})}$ .

## Back to bounding random walks.



**Question:** how is the curve being rescaled?

**Answer:**  $\sqrt{\ln(\#\text{paths})}$ .




## Back to bounding random walks.



**Question:** how is the curve being rescaled?

**Answer:**  $\sqrt{\ln(\#\text{paths})}$ .


## Back to bounding random walks.



**Question:** how is the curve being rescaled?

**Answer:**  $\sqrt{\ln(\#\text{paths})}$ .


## Back to bounding random walks.



**Question:** how is the curve being rescaled?

**Answer:**  $\sqrt{\ln(\#\text{paths})}$ .


## Back to bounding random walks.



**Question:** how is the curve being rescaled?

**Answer:**  $\sqrt{\ln(\#\text{paths})}$ .


## Back to bounding random walks.



**Question:** how is the curve being rescaled?

**Answer:**  $\sqrt{\ln(\#paths)}$ .


## Back to bounding random walks.



**Question:** how is the curve being rescaled?

**Answer:**  $\sqrt{\ln(\#\text{paths})}$ .

## Back to bounding random walks.



**Question:** how is the curve being rescaled?

**Answer:**  $\sqrt{\ln(\#\text{paths})}$ .

## Hoeffding and *one* classifier.

Let  $f$  be a *fixed* classifier. Define random variable

$$Z_i := \mathbb{1} [f(x_i) \neq y_i] .$$

Hoeffding gives: with probability at least  $1 - \delta$ ,

$$\begin{aligned} \Pr[f(X) \neq Y] &= \mathbb{E}(Z_1) \\ &\leq \frac{1}{n} \sum_{i=1}^n Z_i + \sqrt{\frac{\ln(1/\delta)}{2n}} . \\ &\leq \frac{1}{n} \sum_{i=1}^n \mathbb{1}[f(x_i) \neq y_i] + \sqrt{\frac{\ln(1/\delta)}{2n}} . \end{aligned}$$



Hoeffding and *many* classifiers.

This is **exactly** the “many paths” bound.

## Hoeffding and *many* classifiers.

This is **exactly** the “many paths” bound.

Given classifiers  $\mathcal{F}$ , similarly: with probability  $1 - \delta$ , **every**  $f \in \mathcal{F}$  satisfies

$$\Pr[f(X) \neq Y] \leq \frac{1}{n} \sum_{i=1}^n \mathbb{1}[f(x_i) \neq y_i] + \sqrt{\frac{\ln |\mathcal{F}| + \ln(1/\delta)}{2n}}.$$

## Complexity measures.

Given classifiers  $\mathcal{F}$ , with probability  $1 - \delta$ , **every**  $f \in \mathcal{F}$  satisfies

$$\Pr[f(X) \neq Y] \leq \frac{1}{n} \sum_{i=1}^n \mathbb{1}[f(x_i) \neq y_i] + \sqrt{\frac{\ln |\mathcal{F}| + \ln(1/\delta)}{2n}}.$$

## Complexity measures.

Given classifiers  $\mathcal{F}$ , with probability  $1 - \delta$ , **every**  $f \in \mathcal{F}$  satisfies

$$\Pr[f(X) \neq Y] \leq \frac{1}{n} \sum_{i=1}^n \mathbb{1}[f(x_i) \neq y_i] + \sqrt{\frac{\ln |\mathcal{F}| + \ln(1/\delta)}{2n}}.$$

## Generalization.

Given predictors  $\mathcal{F}$ , with probability at least  $1 - \delta$ , each  $f \in \mathcal{F}$  satisfies

$$\text{Risk}(f) \leq \widehat{\text{Risk}}(f) + \tilde{\mathcal{O}} \left( \sqrt{\frac{\text{Complexity}(\mathcal{F}) + \ln(1/\delta)}{n}} \right).$$

where

$$\text{Risk}(f) := \mathbb{E} \ell(f, X, Y) \quad \text{and} \quad \widehat{\text{Risk}}(f) := \frac{1}{n} \sum_{i=1}^n \ell(f, x_i, y_i).$$

## Complexity measures.

Given classifiers  $\mathcal{F}$ , with probability  $1 - \delta$ , **every**  $f \in \mathcal{F}$  satisfies

$$\Pr[f(X) \neq Y] \leq \frac{1}{n} \sum_{i=1}^n \mathbb{1}[f(x_i) \neq y_i] + \sqrt{\frac{\ln |\mathcal{F}| + \ln(1/\delta)}{2n}}.$$

## Generalization.

Given predictors  $\mathcal{F}$ , with probability at least  $1 - \delta$ , each  $f \in \mathcal{F}$  satisfies

$$\text{Risk}(f) \leq \widehat{\text{Risk}}(f) + \tilde{\mathcal{O}} \left( \sqrt{\frac{\text{Complexity}(\mathcal{F}) + \ln(1/\delta)}{n}} \right).$$

where

$$\text{Risk}(f) := \mathbb{E} \ell(f, X, Y) \quad \text{and} \quad \widehat{\text{Risk}}(f) := \frac{1}{n} \sum_{i=1}^n \ell(f, x_i, y_i).$$

**Remark:** holds for all of  $\mathcal{F}$ , in particular for *selected*  $f \in \mathcal{F}$ .

## Complexity measures.

Given classifiers  $\mathcal{F}$ , with probability  $1 - \delta$ , **every**  $f \in \mathcal{F}$  (including choice of an algorithm) satisfies

$$\text{Risk}(f) \leq \widehat{\text{Risk}}(f) + \tilde{O} \left( \sqrt{\frac{\text{Complexity}(\mathcal{F}) + \ln(1/\delta)}{n}} \right).$$

where

$$\text{Risk}(f) := \mathbb{E} \ell(f, X, Y) \quad \text{and} \quad \widehat{\text{Risk}}(f) := \frac{1}{n} \sum_{i=1}^n \ell(f, x_i, y_i).$$

## Complexity measures.

Given classifiers  $\mathcal{F}$ , with probability  $1 - \delta$ , **every**  $f \in \mathcal{F}$  (including choice of an algorithm) satisfies

$$\text{Risk}(f) \leq \widehat{\text{Risk}}(f) + \tilde{O} \left( \sqrt{\frac{\text{Complexity}(\mathcal{F}) + \ln(1/\delta)}{n}} \right).$$


where

$$\text{Risk}(f) := \mathbb{E} \ell(f, X, Y) \quad \text{and} \quad \widehat{\text{Risk}}(f) := \frac{1}{n} \sum_{i=1}^n \ell(f, x_i, y_i).$$

### Example complexity measures.

- ▶ When  $|\mathcal{F}| < \infty$ , can use  $\ln |\mathcal{F}|$ .
- ▶ For classification, can use **VC dimension**.
- ▶ More generally, can use **Rademacher complexity**.

Binomials, random walks, coin tosses, classification.



**Questions so far?**



**A few overfitting/generalization asides.**

Aside: scientific experiments.

- ▶ **Standard scientific setup:** collect some data, try to fit various hypotheses to it.

## Aside: scientific experiments.

- ▶ **Standard scientific setup:** collect some data, try to fit various hypotheses to it.
- ▶ **This is like checking multiple random walks!** The confidence intervals *must* grow with further hypotheses!

## Aside: scientific experiments.

- ▶ **Standard scientific setup:** collect some data, try to fit various hypotheses to it.
- ▶ **This is like checking multiple random walks!** The confidence intervals *must* grow with further hypotheses!
- ▶ This observation is at the core of various “crises” in the application of statistics to science, and give the field of **adaptive data analysis**.

Aside: *covering number* complexities.

We have  $\text{Complexity}(\mathcal{F}) \leq \ln |\mathcal{F}|$ ;  
why not **discretize**  $\mathcal{F}$  and apply?

## Aside: *covering number* complexities.

We have  $\text{Complexity}(\mathcal{F}) \leq \ln |\mathcal{F}|$ ;  
why not **discretize**  $\mathcal{F}$  and apply?

If  $\|w - v\|$  small, then **predictions are similar**, meaning

$$|w^\top x - v^\top x| = |(w - v)^\top x| \leq \|w - v\| \|x\|.$$

## Aside: *covering number* complexities.

We have  $\text{Complexity}(\mathcal{F}) \leq \ln |\mathcal{F}|$ ;  
why not **discretize**  $\mathcal{F}$  and apply?

If  $\|w - v\|$  small, then **predictions are similar**, meaning

$$|w^\top x - v^\top x| = |(w - v)^\top x| \leq \|w - v\| \|x\|.$$

**Consequence:** discretizing

$$\mathcal{F} := \{x \mapsto w^\top x : \|w\| \leq 1\}$$

is same as discretizing  $\{w \in \mathbb{R}^d : \|w\| \leq 1\}$ ,

## Aside: *covering number* complexities.

We have  $\text{Complexity}(\mathcal{F}) \leq \ln |\mathcal{F}|$ ;  
why not **discretize**  $\mathcal{F}$  and apply?

If  $\|w - v\|$  small, then **predictions are similar**, meaning

$$|w^\top x - v^\top x| = |(w - v)^\top x| \leq \|w - v\| \|x\|.$$

**Consequence:** discretizing

$$\mathcal{F} := \{x \mapsto w^\top x : \|w\| \leq 1\}$$

is same as discretizing  $\{w \in \mathbb{R}^d : \|w\| \leq 1\}$ ,  
which needs size  $\mathcal{O}((1/\epsilon)^d)$ , thus

$$\text{Complexity}(\mathcal{F}) \leq \mathcal{O}(d \ln(1/\epsilon)).$$



## Aside: VC dimension.

**VC dimension** gives a complexity measure for classifiers (binary output).

- ▶ **Definition:** the largest data set size (or  $\infty$ ) which this function class (model) can label in all possible ways.

In earlier bounds, can plug in  $\text{Complexity}(\mathcal{F}) \leq \mathcal{O}(\text{VC}(\mathcal{F}) \ln(n))$ .

## Aside: VC dimension.

**VC dimension** gives a complexity measure for classifiers (binary output).

- ▶ **Definition:** the largest data set size (or  $\infty$ ) which this function class (model) can label in all possible ways.

In earlier bounds, can plug in  $\text{Complexity}(\mathcal{F}) \leq \mathcal{O}(\text{VC}(\mathcal{F}) \ln(n))$ .

### Examples.

- ▶ **Linear separators:**  $d$ .
- ▶ **ReLU networks:**  $\tilde{\mathcal{O}}(\#\text{parameters} \cdot \#\text{layers})$ .

## Aside: neural networks.

- ▶ Our generalization bound scales with  $\frac{\text{Complexity}(\mathcal{F})}{n}$ .

## Aside: neural networks.

- ▶ Our generalization bound scales with  $\frac{\text{Complexity}(\mathcal{F})}{n}$ .
- ▶  $\text{Complexity}(\mathcal{F}) \leq \mathcal{O}(\text{VC}(\mathcal{F}) \ln(n))$ , and ReLU networks have  $\tilde{\mathcal{O}}(\#\text{parameters} \cdot \#\text{layers})$ .

## Aside: neural networks.

- ▶ Our generalization bound scales with  $\frac{\text{Complexity}(\mathcal{F})}{n}$ .
- ▶  $\text{Complexity}(\mathcal{F}) \leq \mathcal{O}(\text{VC}(\mathcal{F}) \ln(n))$ , and ReLU networks have  $\tilde{\mathcal{O}}(\#\text{parameters} \cdot \#\text{layers})$ .
- ▶ VC view says: ReLU networks need  $\#\text{parameters} \cdot \#\text{layers} < n$ .

## Aside: neural networks.

- ▶ Our generalization bound scales with  $\frac{\text{Complexity}(\mathcal{F})}{n}$ .
- ▶  $\text{Complexity}(\mathcal{F}) \leq \mathcal{O}(\text{VC}(\mathcal{F}) \ln(n))$ , and ReLU networks have  $\tilde{\mathcal{O}}(\#\text{parameters} \cdot \#\text{layers})$ .
- ▶ VC view says: ReLU networks need  $\#\text{parameters} \cdot \#\text{layers} < n$ .
- ▶ **False in practice** (for further discussion, google “deep learning rethinking generalization”).

## Aside: neural networks.

- ▶ Our generalization bound scales with  $\frac{\text{Complexity}(\mathcal{F})}{n}$ .
- ▶  $\text{Complexity}(\mathcal{F}) \leq \mathcal{O}(\text{VC}(\mathcal{F}) \ln(n))$ , and ReLU networks have  $\tilde{\mathcal{O}}(\#\text{parameters} \cdot \#\text{layers})$ .
- ▶ VC view says: ReLU networks need  $\#\text{parameters} \cdot \#\text{layers} < n$ .
- ▶ **False in practice** (for further discussion, google “deep learning rethinking generalization”).
- ▶ Can use other bounds: “small norm” property(?) of SGD and norm-based generalization.

## Aside: Rademacher complexity.

VC is combinatorial/discrete;

Rademacher is a generalization that allows real-valued predictors.

- ▶ **Definition.** Let  $(\epsilon_1, \dots, \epsilon_n)$  be **random sign** (“Rademacher”) random variables, meaning  $\Pr[\epsilon_i = +1] = \Pr[\epsilon_i = -1] = 1/2$ , and all are independent. Then

$$\text{Rad}(\mathcal{F}) := \sup_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n \epsilon_i f(x_i).$$

- ▶ **Intuition:** how well  $\mathcal{F}$  fits **random signs**.
- ▶ **Remark:** VC was **worst case** sign patterns.
- ▶ **Remark:** can plug in  $\text{Complexity}(\mathcal{F}) \leq n \cdot \text{Rad}(\mathcal{F})^2$ .



## Aside: Rademacher complexity.

VC is combinatorial/discrete;

Rademacher is a generalization that allows real-valued predictors.

- ▶ **Definition.** Let  $(\epsilon_1, \dots, \epsilon_n)$  be **random sign** (“Rademacher”) random variables, meaning  $\Pr[\epsilon_i = +1] = \Pr[\epsilon_i = -1] = 1/2$ , and all are independent. Then

$$\text{Rad}(\mathcal{F}) := \sup_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n \epsilon_i f(x_i).$$

- ▶ **Intuition:** how well  $\mathcal{F}$  fits **random signs**.
- ▶ **Remark:** VC was **worst case** sign patterns.
- ▶ **Remark:** can plug in  $\text{Complexity}(\mathcal{F}) \leq n \cdot \text{Rad}(\mathcal{F})^2$ .

Linear predictors with norm  $\leq R$ : then  $\text{Rad}(\mathcal{F}) \leq R/\sqrt{n}$ .

Neural networks with weight matrices  $(W_1, \dots, W_L)$ :

$$\text{Rad}(\mathcal{F}) = \tilde{\mathcal{O}} \left( (\text{gross stuff}) \cdot \prod_{i=1}^L \sigma_{\max}(W_i) \right).$$

Aside: ridge regression.

## Aside: ridge regression.

Recall the *Ridge Regression Estimator* in matrix/vector form:

$$\hat{w}_\lambda := \arg \min_{w \in \mathbb{R}^d} \frac{1}{2n} \|Xw - y\|_2^2 + \frac{\lambda}{2} \|w\|_2^2.$$

Role of  $\lambda$  in standard learning theory questions:

## Aside: ridge regression.

Recall the *Ridge Regression Estimator* in matrix/vector form:

$$\hat{w}_\lambda := \arg \min_{w \in \mathbb{R}^d} \frac{1}{2n} \|Xw - y\|_2^2 + \frac{\lambda}{2} \|w\|_2^2.$$

Role of  $\lambda$  in standard learning theory questions:

- **Representation:** not just a linear predictor, moreover  $\mathcal{F}_\lambda := \{x \mapsto w^\top x : \|w\| \leq \sqrt{1/\lambda}\}.$

## Aside: ridge regression.

Recall the *Ridge Regression Estimator* in matrix/vector form:

$$\hat{w}_\lambda := \arg \min_{w \in \mathbb{R}^d} \frac{1}{2n} \|Xw - y\|_2^2 + \frac{\lambda}{2} \|w\|_2^2.$$

Role of  $\lambda$  in standard learning theory questions:

- ▶ **Representation:** not just a linear predictor, moreover  $\mathcal{F}_\lambda := \{x \mapsto w^\top x : \|w\| \leq \sqrt{1/\lambda}\}$ .
- ▶ **Optimization:** as discussed in class, to achieve accuracy  $\epsilon$ , gradient descent needs  $\frac{\sigma_{\max}(X) + \lambda}{\sigma_{\min}(X) + \lambda} \ln(1/\epsilon)$  iterations.

## Aside: ridge regression.

Recall the *Ridge Regression Estimator* in matrix/vector form:

$$\hat{w}_\lambda := \arg \min_{w \in \mathbb{R}^d} \frac{1}{2n} \|Xw - y\|_2^2 + \frac{\lambda}{2} \|w\|_2^2.$$

Role of  $\lambda$  in standard learning theory questions:

- ▶ **Representation:** not just a linear predictor, moreover  $\mathcal{F}_\lambda := \{x \mapsto w^\top x : \|w\| \leq \sqrt{1/\lambda}\}$ .
- ▶ **Optimization:** as discussed in class, to achieve accuracy  $\epsilon$ , gradient descent needs  $\frac{\sigma_{\max}(X) + \lambda}{\sigma_{\min}(X) + \lambda} \ln(1/\epsilon)$  iterations.
- ▶ **Generalization:** via **Rademacher complexity** and above representation bound, get  $\text{Complexity}(\mathcal{F}_\lambda) \leq 1/\lambda$ .

Aside: online learning and the perceptron algorithm.

(Details in lecture.)

Summary (of overfitting/generalization).



## Summary (of overfitting/generalization).

- ▶ **Intuition:** predictors/model too flexible  $\implies$  overfit (unless tons of data).
- ▶ **Rigorous form:** we have **generalization bounds**, namely bounds between training and test errors of the form

$$\text{Risk}(f) \leq \widehat{\text{Risk}}(f) + \tilde{O} \left( \sqrt{\frac{\text{Complexity}(\mathcal{F}) + \ln(1/\delta)}{n}} \right),$$

and we gave a few definitions of  $\text{Complexity}(\mathcal{F})$ .