

Lecture 17 — k -means.

Alex Schwing and Matus Telgarsky

March 27, 2018

Announcements.

- ▶ Midterm is being graded right now.
Grades posted tonight or tomorrow.
Midterm handed back Thursday after lecture.
Regrade requests have two weeks.
(Sorry we are late with this!!!)
- ▶ **Homeworks pushed back** one week;
no homework due this week;
no TA office hours this week.

Schedule for today.

- ▶ Clustering basics; k -means objective.
- ▶ k -means algorithms.
- ▶ Applications.
- ▶ Ancillary topics.

Reading: Murphy book, chapter 11.

Clustering basics; unsupervised learning.

Clustering basics; unsupervised learning.

Currently we're doing **unsupervised learning**:

- ▶ Data comes **without supervision/labels**;
- ▶ Task is to **find structure in data**.

Clustering basics; unsupervised learning.

Currently we're doing **unsupervised learning**:

- ▶ Data comes **without supervision/labels**;
- ▶ Task is to **find structure in data**.

Schedule:

- ▶ Last lecture: **PCA**.

$$\arg \min_{\substack{\text{subspaces } L \subseteq \mathbb{R}^d \\ \dim(L)=k}} \frac{1}{n} \sum_{i=1}^n \|x_i - \Pi_L x_i\|^2.$$

- ▶ This lecture: ***k*-means clustering**.
- ▶ Future lectures: GMMs, HMMs, EM, GANs,

Clustering basics.

Clustering.

- ▶ *Partition* data $(x_i)_{i=1}^n$ into *clusters*.
- ▶ Similar data in same cluster;
dissimilar data in different clusters.

Clustering basics.

Clustering.

- ▶ Partition data $(x_i)_{i=1}^n$ into *clusters*.
- ▶ Similar data in same cluster;
dissimilar data in different clusters.
(Good clustering depends on good similarity measure!)

Clustering basics.

Clustering.

- ▶ Partition data $(x_i)_{i=1}^n$ into *clusters*.
- ▶ Similar data in same cluster;
dissimilar data in different clusters.
(Good clustering depends on good similarity measure!)

Exemplar-based clustering.

- ▶ Associate each cluster with an *exemplar/center*.
- ▶ Many applications heavily use this center.
- ▶ Natural objective function:

$$\arg \min_{\mu_1, \dots, \mu_k} \frac{1}{n} \sum_{i=1}^n \text{sim}(x_i, \mu_j).$$

Remarks.

- ▶ $1/n$ often dropped;
here it strengthens analogy to risks/losses.
- ▶ k -means uses $\text{sim}(x_i, \mu_j) = \frac{1}{2} \|x_i - \mu_j\|_2^2$.

The k -means objective: basic properties.

Define the k -means objective as

$$\sum_{i=1}^n \min_j \|x_i - \mu_j\|_2^2.$$

Remarks.

- ▶ (μ_1, \dots, μ_k) are the k means/exemplars/centers.
- ▶ Can treat $\min_j \|x - \mu_j\|_2^2$ as a (nonconvex!) loss.
- ▶ The k -means objective and the k -means *method* (presented shortly) are often conflated.
- ▶ This is an *exemplar*-based, *hard* clustering.
There are many other types of clustering!

The k -means objective: gradients.

Let's take gradient wrt μ_1 .

The k -means objective: gradients.

Let's take gradient wrt μ_1 .

Define $\mu(x_i) \in (\mu_1, \dots, \mu_k)$, closest center to x_i . (Ignore ties.)

$$\begin{aligned}\nabla_{\mu_1} \sum_{i=1}^n \min_j \|x_i - \mu_j\|^2 &= \nabla_{\mu_1} \left(\sum_{\substack{i \in (1, \dots, n) \\ \mu(x_i) = \mu_1}} \|x_i - \mu_1\|^2 + \sum_{\substack{i \in (1, \dots, n) \\ \mu(x_i) \neq \mu_1}} \|x_i - \mu_1\|^2 \right) \\ &= \sum_{\substack{i \in (1, \dots, n) \\ \mu(x_i) = \mu_1}} 2(x_i - \mu_1) + 0.\end{aligned}$$

The k -means objective: gradients.

Let's take gradient wrt μ_1 .

Define $\mu(x_i) \in (\mu_1, \dots, \mu_k)$, closest center to x_i . (Ignore ties.)

$$\begin{aligned}\nabla_{\mu_1} \sum_{i=1}^n \min_j \|x_i - \mu_j\|^2 &= \nabla_{\mu_1} \left(\sum_{\substack{i \in (1, \dots, n) \\ \mu(x_i) = \mu_1}} \|x_i - \mu_1\|^2 + \sum_{\substack{i \in (1, \dots, n) \\ \mu(x_i) \neq \mu_1}} \|x_i - \mu_1\|^2 \right) \\ &= \sum_{\substack{i \in (1, \dots, n) \\ \mu(x_i) = \mu_1}} 2(x_i - \mu_1) + 0.\end{aligned}$$

Remarks.

- ▶ Setting to 0, get $\mu'_j := \text{mean} \left(\{x_i : \mu(x_i) = \mu_j\} \right)$.
- ▶ Can define multiple algs from here (will come back to this).

The k -means objective: alternate form via assignments.

Let's make the assignment of data to centers explicit:

$$\min_{\mu_1, \dots, \mu_k} \sum_{i=1}^n \min_j \|x_i - \mu_j\|_2^2 = \min_{\mu_1, \dots, \mu_k} \min_{\substack{A \in \{0,1\}^{n \times k} \\ A1=1}} \sum_{i=1}^n \sum_{j=1}^k A_{ij} \|x_i - \mu_j\|_2^2.$$

Remarks.

- ▶ $A \in \{0,1\}^{n \times k}$ assigns data points to centers.
It is a **hard assignment**: each x_i gets exactly one μ_j .
- ▶ Natural to consider **soft clustering** $A \in [0,1]^{n \times k}$, $A1 = 1$.
We'll return to this next lecture.

The k -medians objective!

What if we drop the square:

$$\min_{\mu_1, \dots, \mu_k} \min_j \sum_{i=1}^n \|x_i - \mu_j\|_2.$$

- ▶ Setting derivative to 0,

$$\sum_{x_i \in C_j} \frac{x_i - \mu_j}{\|x_i - \mu_j\|} = 0.$$

where $C_j := \{x_i : \mu(x_i) = \mu_j\}$.

- ▶ **Univariate case:** recovers median.
Multivariate case: “geometric” medians.
- ▶ If the square seems weird,
for now just treat it as giving means not medians.

k-means algorithms.

k -means algorithms.

k -means objective

$$\sum_{i=1}^n \min_j \|x_i - \mu_j\|_2^2.$$

k -means algorithms.

k -means objective

$$\sum_{i=1}^n \min_j \|x_i - \mu_j\|_2^2.$$

As before: applying ∇_{μ_l} and setting to zero gives

$$\mu'_l := \frac{1}{|C_l|} \sum_{x_i \in C_l} x_i,$$

where $C_l := \{x_i : \mu(x_i) = \mu_l\}$ are points with μ_l as closest center.

k -means algorithms.

k -means objective

$$\sum_{i=1}^n \min_j \|x_i - \mu_j\|_2^2.$$

As before: applying ∇_{μ_l} and setting to zero gives

$$\mu'_l := \frac{1}{|C_l|} \sum_{x_i \in C_l} x_i,$$

where $C_l := \{x_i : \mu(x_i) = \mu_l\}$ are points with μ_l as closest center.
Let's turn this into an algorithm.

LLoyd's method. (“ k -means algorithm”.)

1. Choose initial clusters (C_1, \dots, C_k) .
2. Repeat until convergence:
 - 2.1 **(Recenter.)** Set $\mu_j := \text{mean}(C_j)$ for $j \in (1, \dots, k)$.
 - 2.2 **(Reassign).** Update $C_j := \{x_i : \mu(x_i) = \mu_j\}$ for $j \in (1, \dots, k)$
(break ties arbitrarily).


Lloyd's method. (“ k -means algorithm”.)

1. Choose initial clusters (C_1, \dots, C_k) .
2. Repeat until convergence:
 - 2.1 (**Recenter.**) Set $\mu_j := \text{mean}(C_j)$ for $j \in (1, \dots, k)$.
 - 2.2 (**Reassign.**) Update $C_j := \{x_i : \mu(x_i) = \mu_j\}$ for $j \in (1, \dots, k)$ (break ties arbitrarily).


Remarks.

- ▶ $\mathcal{O}(nkd)$ per iteration.
- ▶ Initialization discussed shortly.
- ▶ This is **alternating minimization** on cluster assignments and cluster centers.
- ▶ Procedure terminates: each step can't increase cost, and there are finitely many partitions.


Static demo.




Static demo.




Static demo.



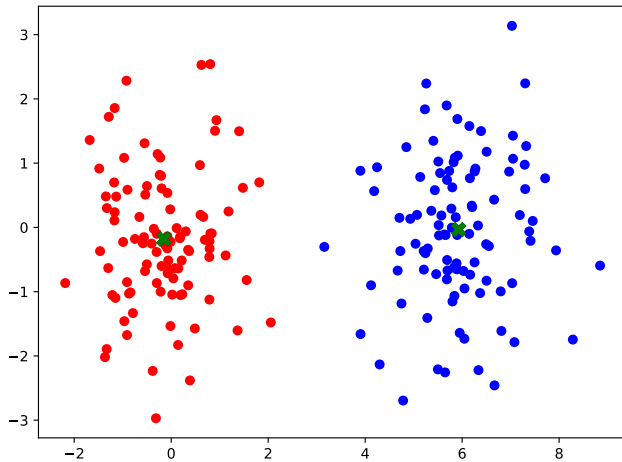
Static demo.



Static demo.



Static demo.



Lloyd's method revisited.

1. Choose initial clusters (C_1, \dots, C_k) .
2. Repeat until convergence:
 - 2.1 **(Recenter.)** Set $\mu_j := \text{mean}(C_j)$ for $j \in (1, \dots, k)$.
 - 2.2 **(Reassign.)** Update $C_j := \{x_i : \mu(x_i) = \mu_j\}$ for $j \in (1, \dots, k)$ (break ties arbitrarily).

Lloyd's method revisited.

1. Choose initial clusters (C_1, \dots, C_k) .
2. Repeat until convergence:
 - 2.1 **(Recenter.)** Set $\mu_j := \text{mean}(C_j)$ for $j \in (1, \dots, k)$.
 - 2.2 **(Reassign.)** Update $C_j := \{x_i : \mu(x_i) = \mu_j\}$ for $j \in (1, \dots, k)$ (break ties arbitrarily).

Let's understand this geometrically.

- ▶ Centers define a **Voronoi diagram/partition**:
for each μ_j , define cell $V_j := \{x \in \mathbb{R}^d : \mu(x) = \mu_j\}$
(break ties arbitrarily).
- ▶ Reassignment leaves assignment consistent with Voronoi cells.
- ▶ Recentering might shift data outside Voronoi cells!

Interactive demo.

Go to `http://mjt.cs.illinois.edu/htv/`.

(Shown in class.)

(This should make the Voronoi cells clear!)

Does the algorithm optimize well?

k -means objective is NP-hard when $d \geq 2$.

- ▶ In practice, Lloyd's method seems to optimize well;
In theory, output can have **unboundedly poor cost**.
(Example given in class: 4 corners of a rectangle.)
- ▶ In practice, method takes few iterations;
in theory: can take $2^{\Omega(\sqrt{n})}$ iterations!
(Examples of this are painful.)

Initialization matters!

- ▶ **Easy choices:**

- ▶ k random points from dataset.
- ▶ Random partition.

- ▶ **Standard choice** (theory and practice):

“ D^2 -sampling”/kmeans++

1. Choose μ_1 uniformly at random from data.
2. for $j \in (2, \dots, k)$:

2.1 Choose $x_i \propto \min_{l < j} \|x_i - \mu_l\|_2^2$.

- ▶ kmeans++ is *randomized furthest-first traversal*;
regular furthest-first fails with outliers.
- ▶ Scikits-learn and Matlab both default to kmeans++.

Applications

Applications

- ▶ The **clusters** found by k -means are useful to *data analysis*: finding groupings that were hard to see.
- ▶ The **exemplars/centers** are also extremely useful!

Application: vector quantization.

Vector quantization with k -means.

- ▶ Let $(x_i)_{i=1}^n$ be given.
- ▶ run k -means to obtain (μ_1, \dots, μ_k) .
- ▶ Replace each $(x_i)_{i=1}^n$ with $(\mu(x_i))_{i=1}^n$.

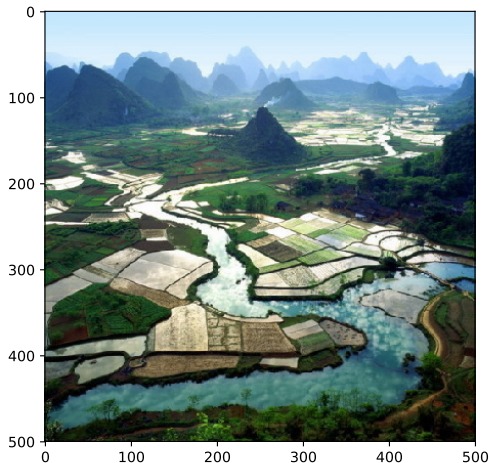
Encoding size reduces from $\mathcal{O}(nd)$ to $\mathcal{O}(kd + n \ln(k))$.

Examples.

- ▶ Audio compression.
- ▶ Image compression.

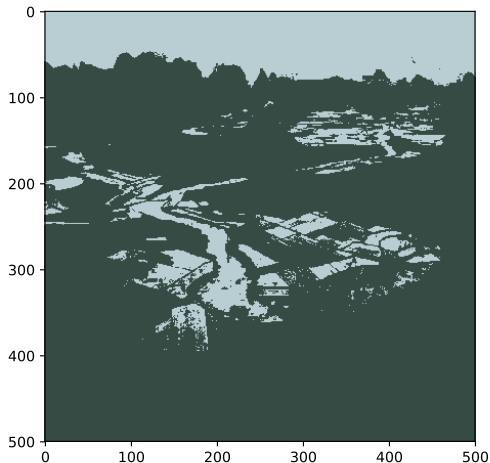
Application: pixel-level quantization.

- ▶ Run k -means on **pixels**.
- ▶ Obtain k exemplars, replace pixels with closest exemplar.



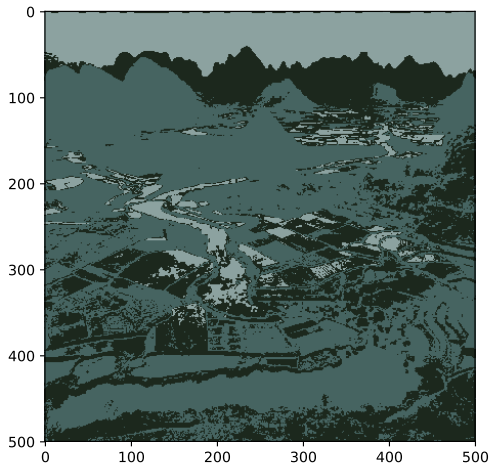
Application: pixel-level quantization.

- ▶ Run k -means on **pixels**.
- ▶ Obtain k exemplars, replace pixels with closest exemplar.



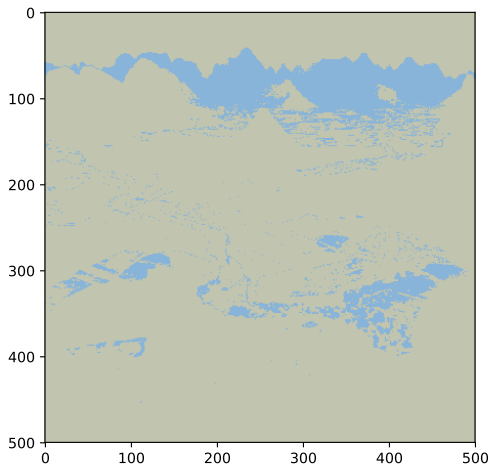
Application: pixel-level quantization.

- ▶ Run k -means on **pixels**.
- ▶ Obtain k exemplars, replace pixels with closest exemplar.




Application: pixel-level quantization.

- ▶ Run k -means on **pixels**.
- ▶ Obtain k exemplars, replace pixels with closest exemplar.




Application: pixel-level quantization.



Looks kindof bad!

Quick fix: use a different *color space*.

Application: pixel-level quantization.



Separating objects in an image is called **segmentation**.

With some tweaks (color space, different clustering method),
can cheaply get a reasonable segmentation.

Application: patch-level quantization.

1. Now $(x_i)_{i=1}^n$ denotes *patches* of *many* images.
2. Obtain exemplars (μ_1, \dots, μ_k) via k -means.
3. Replace image patches with closest exemplar.

Application: feature learning.

1. Start with $(x_i)_{i=1}^n$, where $x_i \in \mathbb{R}^d$.
2. Run k -means, obtain (μ_1, \dots, μ_k) , where $\mu_j \in \mathbb{R}^d$.
3. Replace x_i with $\tilde{x}_i \in \mathbb{R}^k$ where $(\tilde{x}_i)_j := \exp(-\|x_i - \mu_j\|^2)$ (or some other similarity measure).
4. Run whatever ML algorithm (e.g., least squares) on $(\tilde{x}_i)_{i=1}^n$.
(Example in class: the “xor” example we keep mentioning...)

Application: quantizing into “superpixels”.


Earlier quantizations ignore spatial structure.

Cheap fix: add image coordinates to RGB data at each pixel,
and tune the distance metric!

Application: quantizing into “superpixels”.

Earlier quantizations ignore spatial structure.

Cheap fix: add image coordinates to RGB data at each pixel, and tune the distance metric!



Other applications.

Especially after tuning feature encoding and distance metric,
can apply k -means much more broadly.
(E.g., to text.)

Ancillary topics.

Ancillary topics: spherical clusters.

What will k -means do with $k = 3$?



Ancillary topics: spherical clusters.

What will k -means do with $k = 3$?



k -means prefers **spherical clusters**.

Ancillary topics: spherical clusters.

What will k -means do with $k = 3$?




k -means prefers **spherical clusters**.

Changing similarity measure means all clusters still same shape.
Next lecture gives another option.

How to choose k ?


Costs when quantizing pixels of *Guilin*.



Reasonable to choose k at “elbow”;
trades off accuracy and complexity.

How to choose k ?

Costs when quantizing pixels of *Guilin*.



Reasonable to choose k at “elbow”;
trades off accuracy and complexity.

There are other, more complicated ways. (E.g., “bayes information criterion”, “Akaike information criterion”, ...)

Choice of k ; sometimes no good choice.

Which of $k \in \{2, 3\}$ better on following data?



Choice of k ; sometimes no good choice.

Which of $k \in \{2, 3\}$ better on following data?




Perhaps **neither**; want 6. (“Elbow method” works here.)

Choice of k ; hierarchical clustering.

Sometimes *multiple* choices of k make sense.

E.g., Phylogenetic trees have multiple notions of scales.



(Image credit: https://www.researchgate.net/figure/Phylogenetic-tree-based-on-neighbor-joining-analyses-of-a-fig8_260108945.)

The k -means objective: alternate form without exemplars!

Let C_j be the points assigned to μ_j .

$$\sum_{i=1}^n \min_j \|x_i - \mu_j\|_2^2 = \sum_{j=1}^k \frac{1}{2|C_j|} \sum_{a,b \in C_j} \|a - b\|_2^2.$$

This gives a **non-exemplar** way to reason about k -means.

k-means: key points.

k-means: key points.

- ▶ *k*-means is a (hard) clustering method.
- ▶ The objective function is $\min_{\mu_1, \dots, \mu_k} \sum_{i=1}^n \min_j \|x_i - \mu_j\|_2^2$.
- ▶ Remember the standard heuristic (“Lloyd’s method”), it is **alternating minimization** between **assignments** and **centers**.
- ▶ *k*-means finds means/exemplars/centers; these are useful in many applications, e.g., **vector quantization**.

Next time: **Gaussian Mixture Models!**