

Machine Learning

A. G. Schwing & M. Telgarsky

University of Illinois at Urbana-Champaign, 2018

L24: Q-learning

Goals of this lecture

Goals of this lecture

- Extending MDPs

Goals of this lecture

- Extending MDPs
- Getting to know Q-learning

Recap so far: Known MDP

- To compute V^* , Q^* , π^* : use **policy/value iteration or exhaustive**
- To evaluate fixed policy π : use policy evaluation



But what if:

- Transition probabilities and rewards are not known
- No model available (model free RL)

But what if:

- Transition probabilities and rewards are not known
- No model available (model free RL)

Bellman optimality principle:

$$Q^*(s, a) = \sum_{s' \in \mathcal{S}} P(s' \mid s, a) \left[R(s, a, s') + \max_{a' \in \mathcal{A}_{s'}} Q^*(s', a') \right]$$

But what if:

- Transition probabilities and rewards are not known
- No model available (model free RL)

Bellman optimality principle:

$$Q^*(s, a) = \sum_{s' \in \mathcal{S}} P(s' \mid s, a) \left[R(s, a, s') + \max_{a' \in \mathcal{A}_{s'}} Q^*(s', a') \right]$$

Idea:

But what if:

- Transition probabilities and rewards are not known
- No model available (model free RL)

Bellman optimality principle:

$$Q^*(s, a) = \sum_{s' \in \mathcal{S}} P(s' | s, a) \left[R(s, a, s') + \max_{a' \in \mathcal{A}_{s'}} Q^*(s', a') \right]$$

Idea:

- Run a simulator to collect experience tuples/samples (s, a, r, s')

But what if:

- Transition probabilities and rewards are not known
- No model available (model free RL)

Bellman optimality principle:

$$Q^*(s, a) = \sum_{s' \in \mathcal{S}} P(s' \mid s, a) \left[R(s, a, s') + \max_{a' \in \mathcal{A}_{s'}} Q^*(s', a') \right]$$

Idea:

- Run a simulator to collect experience tuples/samples (s, a, r, s')
- Approximate transition probability using samples

Algorithm sketch for Q-learning:

Algorithm sketch for Q-learning:

- Obtain a sample transition (s, a, r, s')

Algorithm sketch for Q-learning:

- Obtain a sample transition (s, a, r, s')
- Obtained sample suggests:

$$Q(s, a) \approx y_{(s,a,r,s')} = r + \max_{a' \in \mathcal{A}_{s'}} Q(s', a')$$

Algorithm sketch for Q-learning:

- Obtain a sample transition (s, a, r, s')
- Obtained sample suggests:

$$Q(s, a) \approx y_{(s,a,r,s')} = r + \max_{a' \in \mathcal{A}_{s'}} Q(s', a')$$

- To account for missing transition probability we keep running average

Algorithm sketch for Q-learning:

- Obtain a sample transition (s, a, r, s')
- Obtained sample suggests:

$$Q(s, a) \approx y_{(s,a,r,s')} = r + \max_{a' \in \mathcal{A}_{s'}} Q(s', a')$$

- To account for missing transition probability we keep running average

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha y_{(s,a,r,s')}$$

Summary:


- Known MDP

- ▶ To compute V^* , Q^* , π^* : use value/policy iteration
- ▶ To evaluate fixed policy π : use policy evaluation



- Unknown MDP: Model free

- ▶ To compute V^* , Q^* , π^* : use Q-learning
- ▶ To evaluate fixed policy π : use value learning

Our current Q-learning works fine in case we have a tabular environment:



Our current Q-learning works fine in case we have a tabular environment:



How to use this idea for playing Atari games?

How to use this idea for playing Atari games?

- What are the actions?

How to use this idea for playing Atari games?


- What are the actions?
- What are the rewards?

How to use this idea for playing Atari games?

- What are the actions?
- What are the rewards?
- What are the states?

How about $Q_{\theta}(s, a)$ as a neural net which takes images as input and produces a number for each of the possible actions?

How about $Q_{\theta}(s, a)$ as a neural net which takes images as input and produces a number for each of the possible actions?



Deep Q-learning algorithm (Deep Q-networks (DQN)):

Given dataset $\mathcal{D} = \{(s_j, a_j, r_j, s_{j+1})\}$:

Deep Q-learning algorithm (Deep Q-networks (DQN)):

Given dataset $\mathcal{D} = \{(s_j, a_j, r_j, s_{j+1})\}$:

- Sample minibatch $\mathcal{B} \subseteq \mathcal{D}$

Deep Q-learning algorithm (Deep Q-networks (DQN)):

Given dataset $\mathcal{D} = \{(s_j, a_j, r_j, s_{j+1})\}$:

- Sample minibatch $\mathcal{B} \subseteq \mathcal{D}$
- Compute target $y_j = r_j + \gamma \max_a Q_{\theta^-}(s_{j+1}, a) \quad \forall j \in \mathcal{B}$

Deep Q-learning algorithm (Deep Q-networks (DQN)):

Given dataset $\mathcal{D} = \{(s_j, a_j, r_j, s_{j+1})\}$:

- Sample minibatch $\mathcal{B} \subseteq \mathcal{D}$
- Compute target $y_j = r_j + \gamma \max_a Q_{\theta^-}(s_{j+1}, a) \quad \forall j \in \mathcal{B}$
- Use stochastic (semi-)gradient descent to optimize w.r.t. parameters θ

$$\min_{\theta} \sum_{(s_j, a_j, r_j, s_{j+1}) \in \mathcal{B}} (Q_{\theta}(s_j, a_j) - y_j)^2$$

Deep Q-learning algorithm (Deep Q-networks (DQN)):


Given dataset $\mathcal{D} = \{(s_j, a_j, r_j, s_{j+1})\}$:

- Sample minibatch $\mathcal{B} \subseteq \mathcal{D}$
- Compute target $y_j = r_j + \gamma \max_a Q_{\theta^-}(s_{j+1}, a) \quad \forall j \in \mathcal{B}$
- Use stochastic (semi-)gradient descent to optimize w.r.t. parameters θ

$$\min_{\theta} \sum_{(s_j, a_j, r_j, s_{j+1}) \in \mathcal{B}} (Q_{\theta}(s_j, a_j) - y_j)^2$$

- Perform ϵ -greedy action and augment \mathcal{D}

Results:



Quiz:

Quiz:

- What differentiates RL from supervised learning?

Quiz:

- What differentiates RL from supervised learning?
- What is a MDP?

Quiz:

- What differentiates RL from supervised learning?
- What is a MDP?
- What to do if no transition probabilities are available?

Important topics of this lecture

Important topics of this lecture

- Getting a feeling for reinforcement learning

Important topics of this lecture

- Getting a feeling for reinforcement learning
- Understanding how to use reinforcement learning

Important topics of this lecture

- Getting a feeling for reinforcement learning
- Understanding how to use reinforcement learning

What's next:

Policy Gradient Methods